

9주차 회귀분석 심화 강의 코드

목차

- 일반화 선형 모형 (GLM)
 - 일반화 선형모형 소개
 - 링크함수 소개
 - 이항 데이터 (로지스틱)
 - 이산형 데이터 (포아송)
- 정규화 회귀분석
 - 정규화란? (Regularization, Penalty)
 - Ridge, Lasso, Elastic 회귀 모형
 - 정규화 회귀의 장점 변수선택
- 비선형 모형 (Non-Linear)
 - 비선형 모형의 필요성 (선형 모형의 한계)
 - 스플라인 회귀 (Spline)
 - 기계학습으로의 확장
- 실제 데이터 분석 하기

```
In [24]: import warnings

# 워닝 메시지 필터링
warnings.filterwarnings("ignore")
```

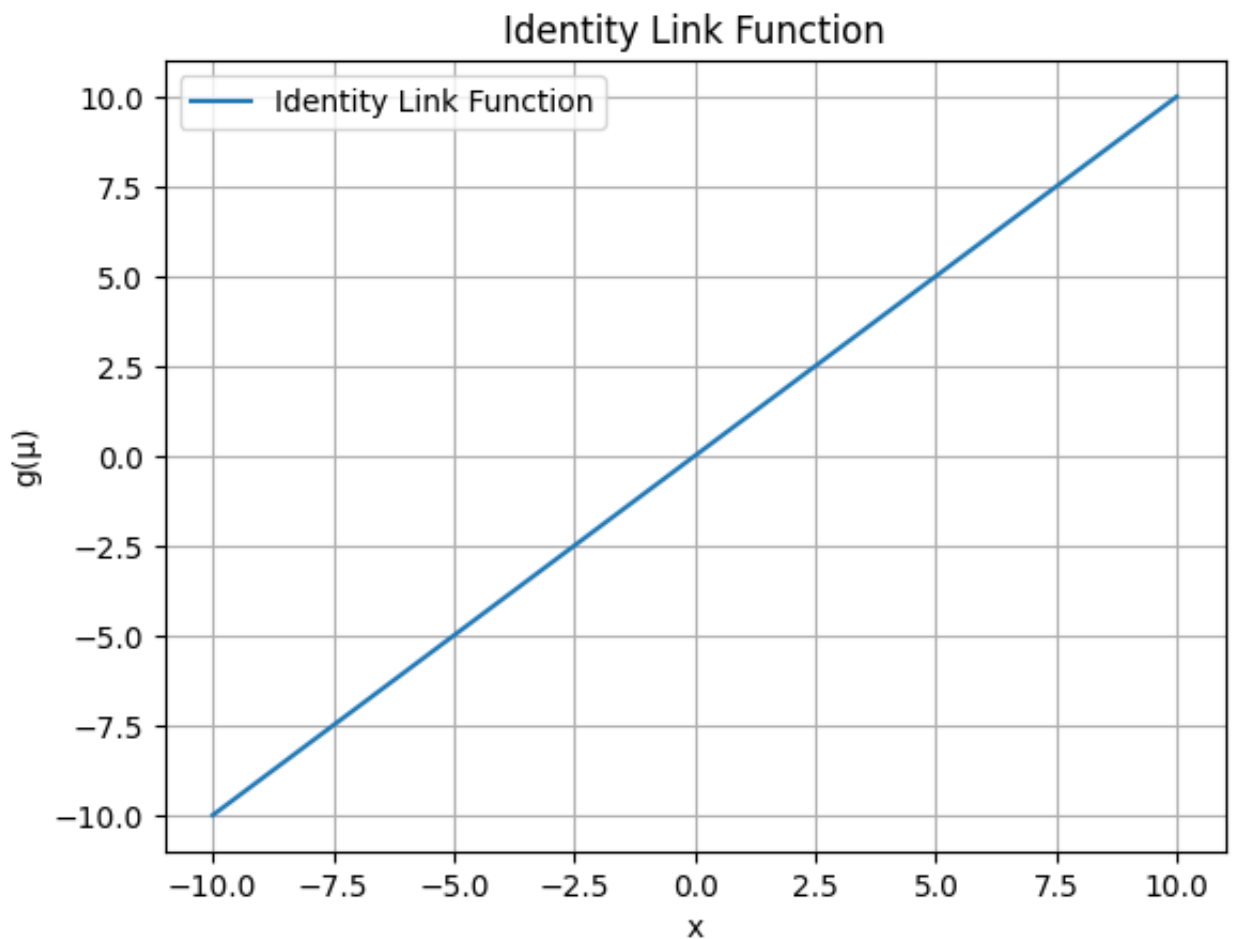
링크함수 소개

- Identity 링크

```
In [2]: import matplotlib.pyplot as plt
import numpy as np

x = np.linspace(-10, 10, 100)
y = x

plt.plot(x, y, label="Identity Link Function")
plt.xlabel("x")
plt.ylabel("g(μ)")
plt.title("Identity Link Function")
plt.legend()
plt.grid(True)
plt.show()
```



- Logit 링크


```
In [8]: import matplotlib.pyplot as plt
import numpy as np

# p는 0에 가깝게부터 1에 가깝게까지의 범위로 설정
p = np.linspace(0.001, 0.999, 400)

# 로짓 함수 계산
logit = np.log(p / (1 - p))

plt.figure(figsize=(10, 6))

# 그래프 그리기
plt.plot(p, logit, label="Logit Function", color='blue')
plt.axhline(0, color='grey', linestyle='--', linewidth=0.5)
plt.axvline(0, color='grey', linestyle='--', linewidth=0.5)
plt.axvline(1, color='grey', linestyle='--', linewidth=0.5)
plt.xlabel("p (Probability)")
plt.ylabel("y (Log Odds)")
plt.title("Logit Function")
plt.legend()
plt.grid(True)
plt.show()
```

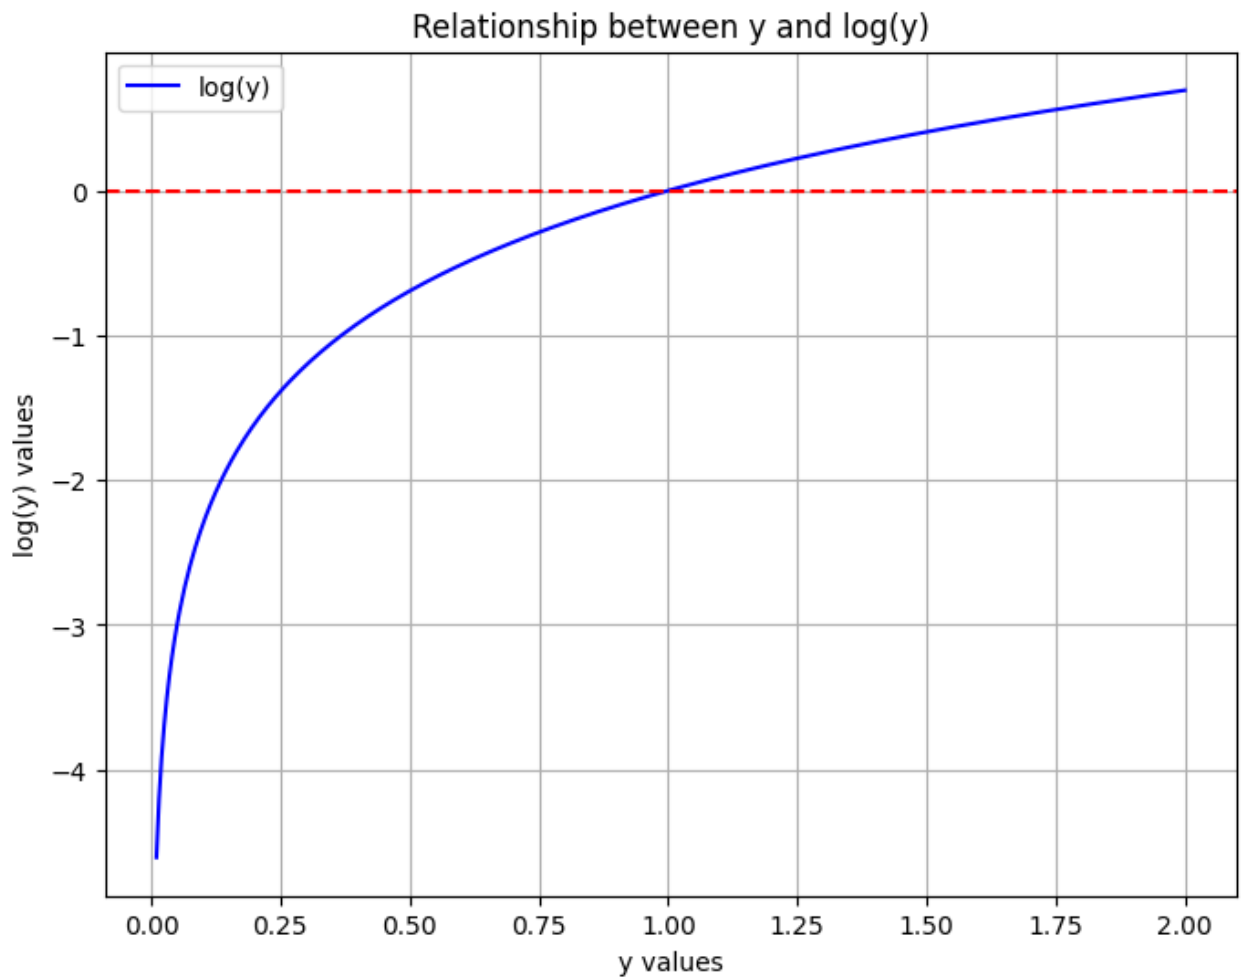


- log 링크

```
In [10]: import numpy as np
import matplotlib.pyplot as plt

# y 값을 0에서 2까지의 범위로 설정
y_values = np.linspace(0.01, 2, 400) # 0에 가까운 값부터 시작해서 log(0)의 정의되지 않음
log_y_values = np.log(y_values)

plt.figure(figsize=(8, 6))
plt.plot(y_values, log_y_values, label='log(y)', color='blue')
plt.axhline(0, color='red', linestyle='--') # y=0 선
plt.xlabel('y values')
plt.ylabel('log(y) values')
plt.title('Relationship between y and log(y)')
plt.legend()
plt.grid(True)
plt.show()
```



일반화 선형모형(GLM)

- 로지스틱

```
In [12]: import numpy as np
import statsmodels.api as sm
import pandas as pd

# 가상의 데이터 생성
np.random.seed(123)
n = 100
x1 = np.random.randn(n)
x2 = np.random.randn(n)
y_continuous = 1 + 2*x1 + 3*x2 + np.random.randn(n)
y = (y_continuous > 0.6).astype(int)

# 데이터프레임으로 변환
data = pd.DataFrame({'x1': x1, 'x2': x2, 'y': y})
data
```

```
Out[12]:
```

	x1	x2	y
0	-1.085631	0.642055	1
1	0.997345	-1.977888	0
2	0.282978	0.712265	1
3	-1.506295	2.598304	1
4	-0.578600	-0.024626	0
...
95	1.031114	-3.231055	0
96	-1.084568	-0.269293	0
97	-1.363472	-0.110851	0
98	0.379401	-0.341262	1
99	-0.379176	-0.217946	0

100 rows × 3 columns

```
In [13]: # 로지스틱 회귀모형 적합
X = sm.add_constant(data[['x1', 'x2']])
model = sm.GLM(data['y'], X, family=sm.families.Binomial())
result = model.fit()

# 결과 출력
print(result.summary())

# 잔차 검정
print("\nResiduals:")
print(result.resid_deviance.head())
```

Generalized Linear Model Regression Results

```

=====
====
Dep. Variable:                y    No. Observations:
100
Model:                        GLM    Df Residuals:
97
Model Family:                Binomial    Df Model:
2
Link Function:                Logit    Scale:                1.
0000
Method:                        IRLS    Log-Likelihood:        -17
.165
Date:                        Fri, 27 Oct 2023    Deviance:                34
.331
Time:                        13:03:15    Pearson chi2:
40.2
No. Iterations:                8    Pseudo R-squ. (CS):        0.
6453
Covariance Type:                nonrobust
=====

```

```

=====
====
              coef      std err          z      P>|z|      [0.025      0.
975]
-----
const          0.8742        0.492        1.777        0.076        -0.090        1
.839
x1             3.8204        0.958        3.989        0.000         1.943        5
.697
x2             6.1775        1.648        3.748        0.000         2.947        9
.408
=====

```

```

Residuals:
0    0.900588
1   -0.032693
2    0.058925
3    0.005307
4   -0.638016
dtype: float64

```

- 포아송 GLM

```

In [16]: # 필요한 라이브러리 임포트
import numpy as np
import pandas as pd
import statsmodels.api as sm

# 데이터 생성
np.random.seed(123) # 재현성을 위한 시드 설정

n_samples = 1000
X = np.linspace(0, 10, n_samples)
# 실제 모델: log(기대값) = intercept + coef * X

```

```

intercept = 1.0
coef = 0.3

# 포아송 분포에서 랜덤하게 샘플 추출
y = np.random.poisson(np.exp(intercept + coef * X))

# 데이터프레임 형태로 변환
data = pd.DataFrame({"X": X, "y": y})

data

```

Out[16]:

	X	y
0	0.00000	2
1	0.01001	5
2	0.02002	3
3	0.03003	0
4	0.04004	2
...
995	9.95996	47
996	9.96997	46
997	9.97998	57
998	9.98999	48
999	10.00000	47

1000 rows × 2 columns

```

In [15]: # 포아송 GLM 모델 적합
exog = sm.add_constant(data["X"]) # 상수항 추가
poisson_model = sm.GLM(data["y"], exog, family=sm.families.Poisson())
result = poisson_model.fit()

# 결과 출력
print(result.summary())

```

Generalized Linear Model Regression Results

```

=====
===
Dep. Variable:                y      No. Observations:
1000
Model:                      GLM      Df Residuals:
998
Model Family:                Poisson  Df Model:
1
Link Function:                Log      Scale:                1.
0000
Method:                      IRLS     Log-Likelihood:        -26
44.7
Date:                        Fri, 27 Oct 2023  Deviance:            10
09.8
Time:                        13:07:54    Pearson chi2:
971.
No. Iterations:                5      Pseudo R-squ. (CS):        1
.000
Covariance Type:                nonrobust
=====
=====
              coef      std err          z      P>|z|      [0.025      0.
975]
-----
const          0.9866      0.024     40.374      0.000      0.939      1
.034
X              0.3011      0.003     93.386      0.000      0.295      0
.307
=====
=====

```

정규화 회귀모형

- 라쏘

```

In [22]: import numpy as np
import statsmodels.api as sm
from sklearn.datasets import make_regression

# 데이터 생성
X, y = make_regression(n_samples=1000, n_features=20, noise=0.1, random_s

# statsmodels를 위한 상수항 추가
X_const = sm.add_constant(X)

pd.DataFrame(X_const)

```


	0	1	2	3	4	5	6	
0	1.0	0.225842	1.551378	-0.107347	0.859695	-0.942963	-1.096625	-1.19
1	1.0	0.110836	-1.454615	0.263888	-1.654510	0.818549	0.482849	0.35
2	1.0	0.458600	-0.081280	-0.698474	0.737528	0.860085	0.275249	0.33
3	1.0	-1.795643	-0.453414	-0.423760	0.155325	0.487775	0.398147	0.73
4	1.0	-1.180626	0.339530	0.328010	-0.224555	0.963951	-1.058450	0.94
...	
995	1.0	1.726964	-0.372833	0.722381	1.024063	-1.760809	0.592527	0.22
996	1.0	0.919229	-1.438278	0.113270	2.062525	1.281016	-1.067533	1.87
997	1.0	-0.512589	1.124777	0.898360	0.906544	-2.301472	0.072252	1.39
998	1.0	-2.968368	-0.929848	0.055208	1.366747	0.427677	0.313143	0.72
999	1.0	-0.487167	2.801373	-1.088635	0.288150	0.321653	0.358848	-0.84

```
# Lasso 회귀 적합 (statsmodels의 GLM 사용)
alpha = 0.5 # 정규화 파라미터
lasso_model = sm.GLM(y, X_const, family=sm.families.Gaussian(), link=sm.g
lasso_results = lasso_model.fit_regularized(method='elastic_net', L1_wt=1

# 결과 출력
print(lasso_results.summary())

# 변수 선택 결과
coef = lasso_results.params
selected_features = np.where(coef != 0)[0]

print("\nSelected features:")
for index in selected_features:
    if index == 0:
        print("Intercept")
    else:
        print(f"Feature {index-1}")
```

=====			
=====			
Dep. Variable:	y	No. Observations:	
1000			
Model:	GLM	Df Residuals:	
990			
Model Family:	Gaussian	Df Model:	
10			
Link Function:	Identity	Scale:	0.01
0377			
Method:	elastic_net	Log-Likelihood:	87
0.15			
Date:	Fri, 27 Oct 2023	Deviance:	10

```

.274
Time:                13:23:26   Pearson chi2:
10.3
No. Iterations:      10   Pseudo R-squ. (CS):      1
.000
Covariance Type:     nonrobust
=====
=====

```

	coef	std err	z	P> z	[0.025	0.
975]						

const	0	0	nan	nan	0	
0						
x1	80.0020	0.003	2.5e+04	0.000	79.996	80
.008						
x2	98.5792	0.003	2.9e+04	0.000	98.573	98
.586						
x3	5.5671	0.003	1748.798	0.000	5.561	5
.573						
x4	0	0	nan	nan	0	
0						
x5	86.4662	0.003	2.67e+04	0.000	86.460	86
.473						
x6	0	0	nan	nan	0	
0						
x7	69.4306	0.003	2.18e+04	0.000	69.424	69
.437						
x8	0	0	nan	nan	0	
0						
x9	0	0	nan	nan	0	
0						
x10	0	0	nan	nan	0	
0						
x11	18.6065	0.003	5928.834	0.000	18.600	18
.613						
x12	39.6357	0.003	1.23e+04	0.000	39.629	39
.642						
x13	0	0	nan	nan	0	
0						
x14	3.1033	0.003	950.599	0.000	3.097	3
.110						
x15	0	0	nan	nan	0	
0						
x16	26.3863	0.003	7891.026	0.000	26.380	26
.393						
x17	0	0	nan	nan	0	
0						
x18	86.8828	0.003	2.66e+04	0.000	86.876	86
.889						
x19	0	0	nan	nan	0	
0						
x20	0	0	nan	nan	0	
0						
=====						
=====						

Selected features:

Feature 0
Feature 1
Feature 2
Feature 4
Feature 6
Feature 10
Feature 11
Feature 13
Feature 15
Feature 17

비선형 회귀모형 소개

- 다항 회귀(Poly Regression)

```
In [32]: import numpy as np
import pandas as pd
import statsmodels.api as sm
from sklearn.datasets import make_regression
from sklearn.preprocessing import PolynomialFeatures

# 데이터 생성 (2개의 설명변수)
X, y = make_regression(n_samples=1000, n_features=2, noise=0.1, random_st

# 다항식으로 데이터 변환
poly_transformer = PolynomialFeatures(degree=2, include_bias=False, inter
X_poly = poly_transformer.fit_transform(X)

#  $X_1^2$ 와  $X_1 \times X_2$  항 제거
X_poly = X_poly[:, [0, 1, 3]] #  $X_1$ ,  $X_2$ ,  $X_2^2$  항만 선택

# 변수명 지정
column_names = ['X1', 'X2', 'X2^2']

# DataFrame으로 변환
X_df = pd.DataFrame(X_poly, columns=column_names)

# statsmodels를 위한 상수항 추가
X_poly_const = sm.add_constant(X_df)

# OLS 회귀 모델 적합
poly_model = sm.OLS(y, X_poly_const).fit()

# 결과 출력
print(poly_model.summary())
```

OLS Regression Results

```

=====
=====
Dep. Variable:          y      R-squared:          1
.000
Model:                  OLS    Adj. R-squared:      1
.000
Method:                  Least Squares    F-statistic:          5.407
e+07
Date:                    Fri, 27 Oct 2023    Prob (F-statistic):
0.00
Time:                    13:30:58    Log-Likelihood:          89
9.63
No. Observations:        1000    AIC:          -1
791.
Df Residuals:            996    BIC:          -1
772.
Df Model:                 3
Covariance Type:         nonrobust
=====
=====

```

	coef	std err	t	P> t	[0.025	0.
975]						
const	0.0002	0.003	0.050	0.960	-0.006	0
.006						
X1	40.7142	0.003	1.25e+04	0.000	40.708	40
.721						
X2	6.6030	0.003	2148.023	0.000	6.597	6
.609						
X2^2	9.439e-05	0.003	0.029	0.977	-0.006	0
.006						

```

=====
=====
Omnibus:                1.906    Durbin-Watson:          2
.013
Prob(Omnibus):           0.386    Jarque-Bera (JB):        1
.791
Skew:                    0.066    Prob(JB):                0
.408
Kurtosis:                3.160    Cond. No.
1.12
=====
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

- 스무딩 스플라인 (Smoothing Spline)

```

In [35]: import numpy as np
from pygam import LinearGAM, s
import matplotlib.pyplot as plt

```

```
# 간단한 비선형 데이터 생성
np.random.seed(42)
X = np.linspace(-10, 10, 1000)
y = np.sin(X) + 0.5 * np.random.normal(size=1000)

# 데이터 형태 변환 (2D array로)
X = X[:, np.newaxis]

# 스무딩 스플라인 적합
gam = LinearGAM(s(0)).fit(X, y)

# 결과 출력
print(gam.summary())

# 예측값 및 신뢰구간 플롯
XX = gam.generate_X_grid(term=0)
plt.plot(XX, gam.predict(XX), 'r-')
plt.plot(XX, gam.predict(XX) + gam.confidence_intervals(XX, width=0.95)[:],
plt.plot(XX, gam.predict(XX) - gam.confidence_intervals(XX, width=0.95)[:],
plt.scatter(X, y, facecolor='gray', edgecolors='none', s=5)
plt.title("Smoothing Spline")
plt.show()
```

LinearGAM

```
=====
Distribution:                      NormalDist Effective DoF:
14.3348
Link Function:                    IdentityLink Log Likelihood:
-1548.9948
Number of Samples:                1000 AIC:
3128.6592
                                     AICc:
3129.1685
                                     GCV:
0.2455
                                     Scale:
0.2392
                                     Pseudo R-Squared:
0.6836
=====
```

Feature Function		Lambda	Rank	EDoF
P > x	Sig. Code			
s(0)		[0.6]	20	14.3
1.11e-16	***			
intercept			1	0.0
9.51e-01				


Significance codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

WARNING: Fitting splines and a linear function to a feature introduces a model identifiability problem
which can cause p-values to appear significant when they are not.

WARNING: p-values calculated in this manner behave correctly for un-penalized models or models with known smoothing parameters, but when smoothing parameters have been estimated, the p-values are typically lower than they should be, meaning that the tests reject the null too readily.

None

Smoothing Spline



In []: